

AI Governance Engineering for Governments EU & Canada

From Policy and Compliance to Operational, Verifiable Control

Ismael Palomo Balañá

CEO & Founder · AstroLife™ Universe

25 September 2026

Abstract

Governments are moving from experimentation with AI towards operational deployment across public services, administrative processes, decision support and critical environments. This paper proposes AI Governance Engineering as an architectural approach for translating governance intent into operational, observable and verifiable capabilities. The focus is not to replace legal or policy frameworks, but to connect them with authority, authorisation, execution controls, evidence, continuous assurance and accountability.

The EU and Canada provide relevant institutional contexts. In the EU, the AI Act establishes a governance and enforcement architecture involving the European Commission's AI Office, national competent authorities and other bodies. In Canada, the federal Directive on Automated Decision-Making and Algorithmic Impact Assessment provide structured requirements for automated decision systems. These frameworks illustrate the need for governance capabilities that remain meaningful across the lifecycle and at the point where systems act.

1. The Public-Sector AI Execution Gap

A policy can define what should happen. An assessment can identify risks. Documentation can establish responsibilities. An audit can reconstruct events. None of these, by itself, guarantees that an AI-enabled system remains within its authorised boundary while operating.

The engineering problem is therefore to connect governance intent with operational state:

Policy → Requirements → Controls → Authorisation → Execution → Evidence → Accountability

For increasingly autonomous systems, the operational chain extends further:

Systems → Models → Agents → Decisions → Authorisation → Execution → Evidence → Accountability

2. EU Governance Context

As of September 2026, the EU AI Act has entered a significant enforcement phase. The European Commission states that enforcement is shared between the AI Office, the European Data Protection

Supervisor and national competent authorities, with certain AI Act enforcement powers applying from 2 August 2026. The governance architecture also includes the European AI Board, Scientific Panel and Advisory Forum. [Ncited0turn0search0Öturn0search1Ö](#)

The European Commission describes the AI Office as supporting implementation and enforcement, including evaluation methodologies, guidance, investigations and corrective measures for relevant providers and systems. [Ncited0turn0search4Ö](#)

For public-sector architecture, this creates an important distinction: regulatory obligations and governance structures define requirements, while operational systems still need mechanisms capable of implementing, observing and evidencing those requirements.

3. Canadian Governance Context

Canada's federal approach includes the Directive on Automated Decision-Making and the Algorithmic Impact Assessment (AIA). The Government of Canada describes the Directive as establishing measures for governing automated decision systems in federal institutions, while the AIA helps assess and mitigate impacts and determine applicable requirements. [Ncited0turn0search5Öturn0search7Ö](#)

Recent Canadian updates have strengthened transparency, senior-level accountability, testing and monitoring, while expanding and refining the AIA and related requirements. [Ncited0turn0search8Ö](#)

This provides a practical illustration of the same architectural challenge: governance requirements must ultimately be connected to system design, decision processes, monitoring, evidence and accountable human roles.

4. From Governance Documentation to Governance Engineering

AI Governance Engineering treats governance as an engineering problem with explicit states, controls, interfaces and evidence. Its purpose is not to turn every governance function into software, nor to replace legal authorities. Rather, it creates an operational layer through which governance requirements can become observable and verifiable.

- Authority — identify the accountable authority and scope of decision rights.
- Authorisation — determine whether a specific action is permitted in context.
- Execution Boundary — constrain what the system can actually invoke or change.
- Runtime Verification — evaluate relevant conditions before and during consequential actions.
- Human Intervention — provide defined escalation, approval, override or interruption paths.
- Evidence — preserve information needed to reconstruct decisions, controls and actions.
- Continuous Reassessment — reassess governance when policy, risk, context or system state changes.

5. Authorisation Is Temporal

A central proposition is that authorisation should not be treated as permanently valid merely because it was granted earlier.

The operational question becomes:

Can I? → Am I authorised? → Am I still authorised now?

A system may have been authorised when a workflow started but encounter a changed risk level, policy condition, identity, data context or system state before execution. Governance engineering therefore needs a temporal dimension: authority and control conditions must remain valid for the action being performed.

6. Continuous Assurance

Continuous assurance extends governance beyond a pre-deployment assessment or post-incident audit. It connects policy intent to observable controls and evidence throughout the lifecycle.

Policy → Control → Implementation → Evidence → Assurance

At the execution boundary:

Authority → Authorisation → Execution → Verification → Evidence → Accountability

The objective is not continuous surveillance for its own sake. It is the ability to determine whether the conditions that justified an action remain valid when the action occurs, and to support reassessment or intervention when they do not.

7. Agentic and Multi-System Public AI

As AI systems become agentic, governance must account for interactions among systems, agents, tools and external services. A single system can remain within its local boundary while a sequence of individually permitted actions produces a materially different aggregate outcome.

A useful operational model is:

Agent → Agent → Tool → System → Outcome

Governance therefore needs to consider delegation, dependency, cumulative effects, cross-system authority and intervention rights, not only the properties of an individual model.

8. Human Control as an Engineered Capability

Human oversight is meaningful only when human authority can be exercised at the points where it matters. This requires defined decision rights, escalation thresholds, approval mechanisms, intervention paths and evidence.

The principle is therefore:

Human control must be engineered, not assumed.

This approach is consistent with the wider international governance direction. The United Nations Global Dialogue on AI Governance identifies transparency, accountability, human oversight and interoperability among its governance priorities. [NcitedÖturn0search9Öturn0search11Ö](#)

9. A Public-Sector Architecture

A government-oriented AI Governance Engineering architecture can be represented as five connected layers:

- Governance Intent — laws, regulations, policies, institutional mandates and principles.
- Governance State — applicable requirements, risk, authority, approvals and current conditions.
- Operational Controls — permissions, policy enforcement, monitoring, intervention and execution constraints.
- Evidence Layer — decisions, authorisations, controls, events, exceptions and outcomes.
- Accountability Layer — responsible roles, review, reconstruction, assurance and remediation.

The architecture should remain interoperable with existing identity, access management, workflow, security, data and execution platforms. AI governance does not need to become the execution engine; it needs to make governance state and authority machine-verifiable where consequential action occurs.

10. Value for Governments

- Trustworthy public services through demonstrable governance controls.
- Operational resilience through continuous monitoring and reassessment.
- Accountability through reconstructable decision and action evidence.
- Human oversight through defined intervention and escalation mechanisms.
- Interoperability through explicit governance interfaces between systems.
- Sovereign capability through governance architectures that can be deployed and controlled within public-sector environments.

The resulting value proposition can be expressed as:

Trusted AI → Resilient Public Services → Sovereign Capability → Verifiable Accountability

11. AstroLife™ Architectural Direction

AstroLife™ Universe is developing this direction through AI Governance Engineering and the AstroLife™ AI Governance OS. The architectural position is to connect governance with operational control rather than treating governance as a separate documentary layer.

Governance → Architecture → Operational Control → Evidence → Accountability

The objective is not to claim exclusivity over an emerging discipline. It is to contribute an architectural perspective focused on operational, verifiable governance for increasingly autonomous and critical AI systems.

12. Research Position

AI Governance Engineering can be understood as the discipline of translating governance principles, requirements and authority into operational capabilities that can be observed, verified, evidenced and reassessed across the AI lifecycle.

For governments, the central question is moving from whether a governance framework exists to whether governance remains effective when an AI-enabled system acts.

Conclusion

The next phase of public-sector AI governance will require more than frameworks, assessments and documentation. It will require architectures in which authority, authorisation, controls, human intervention, evidence and accountability remain connected to system operation.

For Europe, Canada and the wider international community, the engineering challenge is clear: transform governance intent into operational capability without weakening legal authority, democratic accountability or human control.

Govern AI. Build Trust. Preserve Sovereignty.

References and Further Reading

European Commission — AI Act: governance, enforcement and implementation. <https://digital-strategy.ec.europa.eu/en/policies/ai-act-governance-and-enforcement>

European Commission — The enforcement framework of the AI Act. <https://digital-strategy.ec.europa.eu/en/policies/enforcement-ai-act>

European Commission — European AI Office. <https://digital-strategy.ec.europa.eu/en/policies/ai-office>

Government of Canada — Automated decision-making. <https://www.canada.ca/en/government/system/digital-government/digital-government-innovations/responsible-use-ai/automated-decision-making.html>

Government of Canada — Algorithmic Impact Assessment. <https://open.canada.ca/aia-eia-js/>

United Nations — Global Dialogue on AI Governance. <https://www.un.org/global-dialogue-ai-governance/en>

Legal Notice & Confidentiality

© 2026 Ismael Palomo Balañá · AstroLife™. All rights reserved.

This document is an original research and architectural position paper prepared by Ismael Palomo Balañá in his capacity as CEO & Founder of AstroLife™. Unless otherwise indicated, the concepts, terminology, architectural models, diagrams, formulations and original analytical content contained herein are protected by applicable intellectual property and copyright laws.

This document is provided for research, strategic, technical and professional discussion purposes. It does not constitute legal, regulatory, financial, investment, security or compliance advice, nor does it constitute a certification, guarantee of compliance, or assurance that any particular AI system satisfies applicable law or regulatory requirements.

Third-party names, publications, trademarks and referenced materials remain the property of their respective owners. Their inclusion does not imply endorsement, affiliation or certification unless expressly stated.

No part of this document may be reproduced, distributed, modified, commercially exploited or incorporated into another work without appropriate attribution and, where required, prior written authorisation from the rights holder.

Ismael Palomo Balañá
CEO & Founder · AstroLife™ Universe
25 September 2026